

# 1 **GeneFunnel: a mean absolute deviation-based, dispersion-adjusted** 2 **gene set scoring method**

3

4 Emir Turkes<sup>1</sup>, Karen E. Duff<sup>1,2</sup>

5

6 <sup>1</sup>UK Dementia Research Institute at University College London, London, UK

7 <sup>2</sup>Taub Institute, Columbia University Medical Center, New York, NY, USA

8

## 9 **\*Corresponding authors:**

10 Emir Turkes; [emir.turkes.19@ucl.ac.uk](mailto:emir.turkes.19@ucl.ac.uk), Karen E. Duff; [k.duff@ucl.ac.uk](mailto:k.duff@ucl.ac.uk)

11

## 12 **Keywords:**

13 Gene Set Scoring, Gene Set Enrichment, Pathway Analysis, Bioinformatics

14

## 15 **Abstract:**

16 **Background:** Gene set enrichment is central to the interpretation of  
17 transcriptomic and proteomic data. Functional class scoring methods such as  
18 GSVA and ssGSEA provide per-sample pathway activity but can introduce  
19 inter-sample dependence, handle zeros and missing values inconsistently, and  
20 vary widely in computational efficiency and stability. Over-representation  
21 approaches, for example g:Profiler, test predefined hit lists against gene set  
22 catalogues, yet they depend on arbitrary thresholds for differential expression,  
23 are sensitive to list size and background choice, and ignore the magnitude and  
24 evenness of expression across all features.

25 **Results:** We present GeneFunnel, which for each gene set in each sample  
26 computes a dispersion-adjusted sum using a size-aware mean absolute  
27 deviation-derived penalty, discouraging scores driven by a few outlier genes  
28 and favouring more even contributions. The resulting scores lie on an  
29 absolute, sample-independent scale with a baseline at zero, simplifying cross-  
30 dataset comparison and use with methods that assume non-negative inputs. In  
31 simulations probing partial activation, variance-only changes, and subtle  
32 coordinated shifts, GeneFunnel shows high sensitivity with low false-positive  
33 rates when paired with standard statistical testing using limma. In real single-  
34 cell RNA-seq data from Alzheimer's Disease post-mortem brain tissue,  
35 GeneFunnel highlights pathology-relevant processes and down-weights small  
36 sets driven by only a few genes. An Rcpp implementation delivers leading  
37 runtime and memory use when benchmarked against comparable methods.

38 **Conclusions:** GeneFunnel provides fast, interpretable and dispersion-  
39 adjusted per-sample pathway scores that integrate cleanly with common  
40 statistical pipelines and are practical for bulk and single-cell studies. The  
41 software is released as an R package with source code on GitHub  
42 (<https://github.com/eturkes/genefunnel>), alongside a companion web  
43 application for data upload, scoring and results export  
44 (<https://data.duff-lab.org/app/genefunnel-shiny-app>).

## 45 **Introduction:**

46 Gene set enrichment is widely used to interpret high-throughput expression  
47 data by summarising gene-level signals into pathway-level readouts, and has  
48 been reviewed extensively in [1–8]. Approaches fall into three broad families.  
49 Over-representation analysis evaluates predefined hit lists against gene-set  
50 catalogues and is simple to apply, but depends on arbitrary thresholds, is  
51 sensitive to list size and background choice, and discards information about  
52 magnitude and evenness of expression. Rank-based methods popularised by  
53 GSEA [9] assess whether set members concentrate at the extremes of a  
54 ranked list and can detect subtle coordinated shifts, yet they operate at the  
55 group level and are not inherently per-sample. Functional class scoring (FCS)  
56 methods, including GSVA [10] and ssGSEA [11], deliver per-sample scores but  
57 often borrow information across samples through ranking or normalisation,  
58 which can alter scores when the sample set changes and can complicate  
59 comparisons across datasets.

60  
61 In practice, analysts face several recurring problems. First, many procedures  
62 entangle samples during scoring, which weakens the appeal of per-sample  
63 interpretation and undermines reusability of scores across studies. Second,  
64 zeros and missing values are handled inconsistently, despite being common in  
65 RNAseq and proteomics. Zeros in particular are often informative and should  
66 not be silently discarded. Third, very small gene sets can dominate results  
67 when a single highly expressed gene carries most of the signal, whereas large  
68 sets can be favoured simply by size when penalties are not calibrated. Finally,  
69 runtime and memory demands vary greatly between methods, which can limit  
70 routine use on single-cell data.

71  
72 We introduce GeneFunnel, a dispersion-adjusted gene-set scoring method that  
73 retains strict sample independence and produces scores on an absolute scale  
74 with a natural zero baseline. For each gene set in each sample, GeneFunnel  
75 computes the set sum and subtracts a size-aware mean absolute deviation  
76 (MAD) penalty. The score rises with total activity but is reduced when  
77 contributions are uneven, so sets with broadly shared signal score higher than  
78 those dominated by a few genes. Zero counts are treated as measurements  
79 and retained, while missing values are not permitted at scoring; when a gene  
80 set includes features absent from the matrix, the set is restricted to the  
81 features present.

82  
83 GeneFunnel is designed to be intuitive to reason with. Scores increase with  
84 total activity, but only to the extent that contributions are broadly shared  
85 within the set. The size-aware penalty is stronger for smaller sets, which  
86 controls spurious hits from two or three gene terms, and approaches a  
87 constant for large sets. The scoring rule yields non-negative outputs with zero  
88 as the limit under maximal unevenness, which facilitates downstream  
89 transformations and distance measures that assume non-negative inputs. A

performant Rcpp implementation enables routine use on bulk and single-cell matrices in RNAseq, proteomics, or other assays with comparable data types.

### **Implementation:**

The functional class scoring algorithm we propose, GeneFunnel, is straightforward: it is similar to subtracting the mean absolute deviation from the sum of values in a gene set. GeneFunnel iterates through each gene set for each sample, introducing no dependency between samples or between gene sets. For a given gene set in a given sample, we first subset the sample's features to those present in the set and compute the set sum and the set mean. We then take, for each feature in the set, the absolute difference between its expression and the set mean, and sum these absolute deviations. A size-aware scaling factor is applied to this deviation term, defined as the size of the gene set divided by twice the residual gene set size (1 minus the gene set size). Finally, the scaled deviation is subtracted from the set sum to yield the score. This procedure is repeated for all gene sets within the sample and then across all samples, producing a gene-set-by-sample score matrix that mirrors the shape of the original data (Figure 1A). The algorithm is expressed in mathematical notation in Figure 1B, and an excerpt of the Rcpp (C++ interface for R) implementation is shown in Figure 1C.

At the core of GeneFunnel's scoring method is its use of both the sum and deviation of feature expression levels within a gene set. By first summing expression values, the method captures the overall activity level of a pathway, akin to approaches that rely on simple averaging or summation. However, instead of assuming that all features contribute equally, GeneFunnel then computes deviance scores for each feature, measuring how much each feature's expression deviates from the mean expression of the set. This deviation-aware component ensures that pathways with highly variable expression across member features are penalised, preventing scenarios where a small number of highly expressed features dominate the enrichment score.

### **Proof of Non-negative Scores:**

A fundamental requirement for GeneFunnel is that its scores remain non-negative, ensuring compatibility with common downstream analyses in functional genomics such as dimensionality reduction, normalisation, and differential expression analysis. The method is designed such that the minimum possible score is zero, which occurs in two biologically interpretable cases: when all features in the set have zero expression or when the set exhibits maximal internal deviation, meaning that the expression values are so dispersed that the deviation term fully offsets the total summed expression (i.e. the case when a single value is non-zero). Proving that GeneFunnel always produces non-negative scores formally validates that it is a proper transformation of gene expression data, ensuring interpretability and compatibility with standard computational workflows.

135 **Theorem:** GeneFunnel Scores Cannot be Negative

136 Let  $X_{i,j}$  be the expression level of feature  $i$  in sample  $j$ , and let  $G_k$  be a  
137 predefined gene set containing  $|G_k|$  features. The GeneFunnel score for gene  
138 set  $G_k$  in sample  $j$  is given by:

$$\text{score}_{k,j} = \sum_{i \in G_k} X_{i,j} - \left( \frac{|G_k|}{2(|G_k| - 1)} \sum_{i \in G_k} |X_{i,j} - \bar{X}_{G_k,j}| \right)$$

139  
140 Then, for all  $k$  and  $j$ :

$$\text{score}_{k,j} \geq 0.$$

142

143 **Proof:**

144 We begin by expanding the sum of values in the gene set (found left-hand-side  
145 or LHS of the parenthesis), where in a general case, the sum of values is equal  
146 to the mean of values times the number of values:

$$\sum_{i \in G_k} X_{i,j} = |G_k| \bar{X}_{G_k,j}$$

147  
148 Substituting this into the scoring equation, located LHS of the parenthesis, we  
149 obtain:

$$\text{score}_{k,j} = |G_k| \bar{X}_{G_k,j} - \left( \frac{|G_k|}{2(|G_k| - 1)} \sum_{i \in G_k} |X_{i,j} - \bar{X}_{G_k,j}| \right)$$

150  
151 Factoring out  $|G_k|$  from that substitution, and from within the parenthesis,  
152 simplifies the score to:

$$\text{score}_{k,j} = |G_k| \left( \bar{X}_{G_k,j} - \frac{1}{2(|G_k| - 1)} \sum_{i \in G_k} |X_{i,j} - \bar{X}_{G_k,j}| \right)$$

153  
154 Looking within the parenthesis, we form the following inequality, stating that  
155 the mean of values is always greater than the sum of absolute deviances from  
156 the mean multiplied by the scaling factor:

$$\bar{X}_{G_k,j} \geq \frac{1}{2(|G_k| - 1)} \sum_{i \in G_k} |X_{i,j} - \bar{X}_{G_k,j}|$$

157  
158 Note that omitting the scaling factor from the RHS yields the equation for  
159 Mean Absolute Deviation (MAD):

$$\frac{1}{|G_k|} \sum_{i \in G_k} |X_{i,j} - \bar{X}_{G_k,j}|$$

160  
161 As shown in (Aghili-Ashtiani, 2021):

$$\text{MAD} = \frac{1}{n} \sum_{i=1}^n |x_i - \bar{x}|$$

162  
163 Where  $x_1, x_2, \dots, x_n \in \mathbb{R}$  is a set of real numbers.

164 We therefore reformulate the problem as follows, noting however that the  
165 inequality that the mean of values as always being greater than or equal to the  
166 MAD does not hold:

$$\bar{x} \not\geq \frac{1}{n} \sum_{i=1}^n |x_i - \bar{x}|$$

167  
168 In fact, in maximally deviating sets, where there is only a single non-zero  
169 value, the ratio of the MAD to the mean approaches 2 with increasing set size.

170 Let's assume a vector  $x = [a, 0, 0, \dots, 0]$  of length  $n$ , where only the first value is  
171 non-zero. Then:

172  $\bar{x} = \frac{a}{n}$

173 and:

174  $\text{MAD} = \frac{2a(n-1)}{n^2}$

175 The ratio between the MAD and mean can then be written as:

176  $\frac{\text{MAD}}{\bar{x}} = \frac{2a(n-1)/n^2}{a/n} = \frac{2(n-1)}{n}$

177 Finally, as set size approaches infinity:

178  $\lim_{n \rightarrow \infty} \frac{2(n-1)}{n} = \lim_{n \rightarrow \infty} \left(2 - \frac{2}{n}\right) = 2$

179 Therefore:

180  $\bar{x} \not\geq \frac{1}{n} \sum_{i=1}^n |x_i - \bar{x}|$

181 These findings helped influence discovery of the appropriate GeneFunnel  
182 scaling factor. With the scaling factor applied, the above evaluates as follows:  
183

184  $\text{MAD}_{\text{scaled}} = \frac{1}{2(n-1)} \sum_{i=1}^n |x_i - \bar{x}|$

185  $\text{MAD}_{\text{scaled}} = \frac{1}{2(n-1)} \cdot \frac{2a(n-1)}{n} = \frac{a}{n}$

186  $\frac{\text{MAD}_{\text{scaled}}}{\bar{x}} = \frac{a/n}{a/n} = 1$

188  
189 Therefore:

190  $\bar{x} \geq \frac{1}{2(n-1)} \sum_{i=1}^n |x_i - \bar{x}|$

191 Substituting in the definitions for GeneFunnel:

192  $\bar{X}_{G_k, j} \geq \frac{1}{2(|G_k| - 1)} \sum_{i \in G_k} |X_{i, j} - \bar{X}_{G_k, j}|$

193 We show that the inequality is satisfied for GeneFunnel scores, and conclude  
194 that the scaling factor combined with MAD is necessary to ensure that for all  $k$   
195 and  $j$ :

196  $\text{score}_{k, j} \geq 0.$

197  
198 Thus GeneFunnel always produces non-negative scores.  
199  
200  
201  
202  
203  
204  
205

## Results:

### Exploration of GeneFunnel Properties:

To thoroughly evaluate the behaviour of GeneFunnel and understand its scoring properties, we conducted an exploration of its theoretical and practical characteristics before benchmarking it against existing methods. To facilitate this process, we developed a Shiny web application (<https://data.duff-lab.org/app/genefunnel-benchmarks-viewer>), which provides an interactive interface for investigating how GeneFunnel responds to different input scenarios.

A key component of this exploration involved constructing a hypothetical gene by sample matrix to simulate different patterns of gene expression (Figure 2). This synthetic dataset allowed for precise control over the relationships between genes, enabling a systematic examination of how GeneFunnel assigns scores under various conditions. Within the web app, users can interactively modify values within this matrix, effectively simulating different gene expression profiles. Each change is processed in real time, with GeneFunnel recomputing scores for all gene sets dynamically. The results are displayed as a heatmap of the gene set by sample matrix, providing immediate visual feedback on how alterations in individual genes affect pathway-level enrichment scores. This interactive approach not only aids in validating theoretical expectations, such as the behaviour of GeneFunnel under extreme cases, but also helps intuitively illustrate how the method differs from traditional functional class scoring approaches.

The selection of values for the originating gene by sample matrix was very deliberate, to try to cover the broad range of situations GeneFunnel was designed to excel in, within a minimal example. Starting with the first column of Figures 2A and B, with Sample 1, it can be seen that all of the values hover around the arbitrary expression value of 50. These values also include a small degree of noise or jitter. Upon examination of Figure 2C, the metrics below the Sample 1 cell confirm these properties. The mean is precisely 50, and with 10 values, this also results in a sum of 500. The small amount of deviance between features is also captured, which when subtracted from the mean results in a final value of 483.

The values in Sample 2 were specifically selected to contrast with Sample 1. Examining the original values, it is clear that this column contains much more feature deviance, with values above 100 and several values recorded as 0. With the inclusion of zero values, Sample 2 was designed to have the same sum and mean as Sample 1. However, the large total deviance of 298 brings a significant penalty to the final score, dropping it from 500 to just 202. This is in contrast to Sample 1, which has a final score of 483. Sample 3 meanwhile confirms that when all features are equal in value, the total deviance is zero. As the values in this sample now centre around 100 rather than 50, the mean

251 is now 50 while the sum is 1,000. With a lack of feature deviance, this results  
252 in simply an enrichment score equal to the sum.

253  
254 Sample 4 differs substantially from the others in that there are NA values in  
255 the original matrix. As a result, the gene set in Sample 4 is actually treated as  
256 a gene set with a size of 5 rather than 10. This changes the scaling factor.

257 Whereas the other samples have a scaling factor of  $\frac{10}{18}$ , the scaling factor for  
258 Sample 4 changes to  $\frac{5}{8}$ . This would normally have an effect on the deviance  
259 score, though in this example, there is no deviance between features to begin  
260 with, so it remains zero. However, what is noteworthy is that the mean of  
261 Sample 4 is equal to Sample 3, but the final enrichment score and sum is 500  
262 rather than 1,000. This indicates GeneFunnel's preference for scoring larger  
263 gene sets higher, with the argument being that an enriched larger gene set is  
264 more likely to be of biological interest than smaller ones.

265  
266 The final column, Sample 5, showcases a situation where all features exhibit  
267 maximal deviance, producing a score of zero as supported by the proof in  
268 preceding sections. Containing a single non-zero value, the sum is fully  
269 cancelled out by an equivalent deviance penalty. This would be the case in all  
270 gene set sizes containing a single non-zero value. As the proportion of non-  
271 zero values increase, the enrichment score gradually increases until an  
272 equilibrium where half of values are non-zero. Assuming no additional  
273 deviance, the final enrichment score in such case would be half of the sum.

274  
275 **Comparing Properties With Other Functional Class Scoring Methods:**

276 We next aimed to build off the exploratory approach in the preceding section  
277 and apply it to several other functional class scoring methods. We elected to  
278 compare GeneFunnel with other functional class scoring methods that  
279 transform expression data into a gene-set-by-sample matrix: GSVA (testing  
280 both Poisson and Gaussian kernels) [10], ssGSEA [11], PLAGE [12], and the Z-  
281 score method available from the GSVA R package. Like the last analysis, the  
282 results are wholly contained in the web app in the next tab section. The first  
283 series of explorations again focus on a hypothetical gene by sample matrix,  
284 constructed similarly as the first with slight modifications (Figure 3A). After  
285 running each of the models, the results are condensed into enrichment  
286 heatmaps (Figure 3B).

287  
288 All methods were run as recommended by their authors for all benchmarking.  
289 Importantly, all data input into GSVA Gaussian, ssGSEA, PLAGE, and Z-score  
290 underwent a  $\log_2 + 1$  transformation, with GSVA Poisson (and GeneFunnel)  
291 being the only methods taking the raw data. The minimum set size was also  
292 set to 2 for all methods. Finally, the normalisation step in ssGSEA was turned

293 off, as the method is no longer a single-sample method with it applied. All  
294 methods were ran with parallel processing through BiocParallel.

295  
296 Beginning with Sample 1, as the most basic test, both Gene Set X and Y  
297 should be more-or-less similar, as the data contained in each are nearly  
298 identical. GeneFunnel produces scores that reflect this, with 237.75 and  
299 244.50 for Gene Set X and Y, respectively. All other methods show noticeable  
300 and generally large differences between them, especially with GSVA. GSVA in  
301 particular attempts to distribute its output along a range of -1 and 1, similar to  
302 a Z-score, making it the most inappropriate for assessing just a few gene sets.  
303 While this dataset is indeed a contrived example, it is not inconceivable to be  
304 interested in only scoring a few select gene sets in a real-world situation.  
305 GSVA was however, the only method other than GeneFunnel to attribute the  
306 highest score to Gene Set Z, the largest gene set encompassing all features in  
307 the test dataset.

308  
309 In Sample 2, the expectation was to see generally lower scores than in Sample  
310 1, while following the same pattern of Gene Set X and Y being comparable,  
311 and Gene Set Z having the largest scores. GeneFunnel fulfilled this criteria,  
312 while all others failed. Most of the methods showed similar patterning as in  
313 Sample 1, while Z-score appeared to similarly score Gene Set Y and Gene Set  
314 Z (the largest gene set) this time, which could not be explained.

315  
316 Sample 3 is the most straightforward of the samples. With no deviance  
317 between features at all, Gene Set X and Y should be identical, with Gene Set Z  
318 at least being identical or larger. This time, there were two methods that could  
319 be considered comparable to the output seen in GeneFunnel. PLAGE showed  
320 very sensible results in that all gene sets of Sample 3 were the largest scoring  
321 sets in the whole dataset. Furthermore, the scores for Gene Set X and Y are  
322 quite similar (1.643413 vs. 1.594036), though not precisely identical like  
323 GeneFunnel. While Sample 3 Gene Set Z is indeed the highest score in the  
324 entire dataset for GeneFunnel as well, Gene Sets X and Y are more similar to  
325 Gene Set Z of Sample 1. The other method that demonstrated sensible  
326 performance in Sample 3 was GSVA Gaussian, as Gene Set X and Y are more  
327 similar to one another compared to those sets in other samples. Gene Set Z  
328 also received the highest score, though the methodology of GSVA makes it so  
329 that there is no difference in the score of Gene Set Z between samples; it  
330 converges towards 1 in all cases.

331  
332 Finally, it was expected for Sample 4 to produce scores that increase in value  
333 slightly from Gene Set X, to Gene Set Y, to Gene Set Z, as the proportion of  
334 zeros to non-zero values decrease. GeneFunnel reflected this, although the  
335 difference between Gene Set Y and Z were small and hard to see on the  
336 heatmap (50 vs. 66.66). The only other method that had the correct trend was  
337 GSVA Gaussian, however, like other samples, the gap between Gene Set Z



338 compared to the other gene sets is extreme. While PLAGE didn't show the  
339 complete expected pattern (Gene Set Z was the lowest scoring), it did  
340 correctly show Gene Set X as less enriched than Gene Set Y, which is an  
341 undebatable expectation. Furthermore, as a whole, the values in Sample 4 are  
342 the lowest in the entire dataset, which should also be expected.

343  
344 In conclusion, at least in this contrived scenario, aside from GeneFunnel, all of  
345 the tested methods performed in ways that were hard to reason with. While it  
346 appears to be the case that none of these methods were constructed to work  
347 with such small test cases, it is still a significant drawback. After all, a very  
348 useful approach for exploratory work into understanding how a method works  
349 and interacts with changing parameters are through small, controlled  
350 experiments like these. Furthermore, some real-world datasets, particularly in  
351 proteomics, can be of small size, and it is unclear of what size dataset is  
352 Incomparability with such scenarios bring about major limitations to the  
353 adoption of these methods. Outside of this, not every real-world experiment is  
354 high-throughput especially when working with emerging technology such as  
355 spatial omics. It is important for bioinformatic methods to be robust to a range  
356 of dataset sizes and it can be demonstrated here that at least within small  
357 datasets, GeneFunnel performs sensibly.

358  
359 **Benchmarking Against Other Methods in Synthetic Data:**

360 To test whether the small-panel results were an artifact of the setup rather  
361 than the methods, another synthetic benchmark was built comparing two  
362 groups on a large gene catalog and with formal statistical testing alongside a  
363 broader mix of approaches, including approaches that cover the main families  
364 of gene-set inference: ORA, camera, fgsea, and GSVA/ssGSEA. ORA (over-  
365 representation analysis) takes the final list of differentially expressed genes  
366 and asks, via an enrichment test against the background gene catalogue,  
367 whether each set contains more hits than expected by chance; this is the  
368 generalised approach taken by the popular g:Profiler [13], but this  
369 implementation permits an arbitrary set catalogue and background, which is  
370 necessary for synthetic benchmarks. camera, from limma [14, 15], is a  
371 competitive test that fits a linear model per gene and then evaluates whether  
372 genes in a set show stronger differential expression than genes outside the  
373 set. fgsea is an R implementation of GSEA [9], which operates on a ranked list  
374 of genes and computes an enrichment score that reflects whether set  
375 members concentrate near the top or bottom of the ranking [16]. GSVA and  
376 ssGSEA are among the functional class scoring methods used in the prior  
377 benchmark, computing a per-sample score for each set without using group  
378 labels, similar to GeneFunnel. For these, as well as GeneFunnel, a stock  
379 limma-trend pipeline was applied to test for differential enrichment between  
380 groups. This mixture of methods allow for a comparison of hit-list enrichment  
381 (ORA), model-based competitive testing (camera), rank-based enrichment

(fgsea), functional class scoring (GSVA and ssGSEA), and the proposed functional class scoring method (GeneFunnel) under one evaluation protocol.

The simulation uses a 20,000 gene matrix partitioned into 1,000 non-overlapping gene sets (20 genes per set) and 10 samples (5 in group A, 5 in group B). In each experiment, 50 sets are designated signal and the remaining 950 null. Counts are drawn from a negative-binomial model with realistic library-size variation, then normalised with edgeR TMM before set-level scoring and testing across the gene matrix. Signals were injected under three patterns that isolate different behaviours of gene set enrichment methods and expose different dynamic ranges of gene set activity. In “spike”, only half of genes in a signal set is perturbed, but strongly, while the remainder is left untouched, which probes a method’s tolerance to partial activation and within-set heterogeneity (Figure 4A). In “variance”, the set mean of the signal set is preserved while the within-set dispersion is deliberately reduced in one group, testing the ability of the methods in assessing within-set consistency (Figure 4B). In “coordinated”, a small but consistent log fold change is applied to all genes in a signal set in one group, producing the most classic example of gene set enrichment but within a small dynamic range so as to stress the sensitivity of each method (Figure 4C). Each setting was then evaluated at FDR 0.05, using statistical testing intrinsic to each method or limma-trend otherwise, recording sensitivity, specificity, precision and other common benchmarks. In contrast with the last simulation, which functions as an exploration of functional class scoring properties under precisely defined but ultimately unrealistic scenarios, this simulation study intends to more comprehensively cover the various approaches to gene set enrichment in a dataset with realistic properties and signal structures. It furthermore provides a clearer picture of where GeneFunnel’s design, which rewards both signal magnitude and within-set consistency, confers advantages or exposes limitations in practical use.

The tables in Figure 5 summarise method performance at the threshold of FDR (BH adjusted p-value) 0.05 for each perturbation paradigm. With 50 signal sets and 950 null sets per experiment, TP (true positive) counts signal sets correctly detected, FN (false negative) the missed signal sets, TN (true negative) the correctly rejected null sets, and FP (false positive) the null sets falsely called. From these we report sensitivity ( $TP/P$ ), specificity ( $TN/N$ ), precision ( $TP/(TP+FP)$ ), accuracy ( $(TP+TN)/(P+N)$ ), and F1 (the harmonic mean of precision and sensitivity). We also report the average FDR for the signal sets, defined as the mean BH adjusted p-value across all 50 signal sets for each paradigm. Higher is better for all rates except average FDR, where lower is better.

In the “spike” paradigm, where only half of the genes in a signal set are perturbed, though robustly, all methods perform well except ORA. GeneFunnel

427 and camera achieve perfect detection (50/50 true positives with no false  
428 positives). The rank-based and unsupervised scoring approaches are close to  
429 this ceiling; fgsea detects all 50 signal sets with two false positives, GSVA  
430 recovers 45 of 50 with one false positive, and ssGSEA detects all 50 with three  
431 false positives. ORA remains highly specific but has low sensitivity (15/50),  
432 likely because partial activation leaves too few genes surpassing the  
433 differential expression threshold to trigger over-representation at the set  
434 level. Although the 50 signal sets do not overlap any other sets, several  
435 methods are sensitive to the overall distribution of gene-level statistics or  
436 ranks. The robust perturbation of the signal sets may have shifted this  
437 background slightly, resulting in false positives for those methods.  
438 GeneFunnel is only susceptible to this issue at the statistical testing stage, i.e.  
439 limma, as the scoring mechanism itself operates on each gene set and sample  
440 in isolation.

441  
442 In the “variance” paradigm, where the mean is preserved and only within-set  
443 dispersion is altered, procedures that target location differences lose power.  
444 GeneFunnel retains the highest sensitivity because its scoring emphasises  
445 within-set consistency as well as magnitude, allowing reduced variability to  
446 register as a stronger, more coherent pattern despite the lack of mean change.  
447 Camera and ORA identify a smaller fraction of signal sets, and the rank-based  
448 and unsupervised scoring methods detect none at the chosen threshold.  
449 Though no other method claims to measure inter-gene variance, individual  
450 changes to gene counts to reduce inter-gene variance pushes some genes past  
451 the significance threshold for regular differential expression testing. It is  
452 likely that when several such genes occur in the same set, methods that  
453 aggregate gene-level evidence, such as camera or ORA, can incidentally  
454 report enrichment despite not explicitly including criteria for within-set  
455 consistency. Across methods, average FDRs are higher than in “spike”,  
456 reflecting weaker evidence when the signal resides in dispersion rather than  
457 in the mean.

458  
459 In the “coordinated” paradigm, where a very small (0.25 logFC with 0.1  
460 standard deviation), but consistent log-fold change is applied to all members  
461 of each signal set, GeneFunnel again achieves the best combination of  
462 sensitivity and F1. GSVA is second, in line with its design to capture  
463 coordinated per-sample shifts, and camera detects fewer sets at this subtle  
464 effect size. fgsea and ORA do not register signal sets at all here, indicating  
465 that the per-gene effects are too small to accumulate sufficient ranked-list or  
466 hit-list evidence at an FDR of 0.05. This paradigm is the most standard  
467 formulation of gene set enrichment and serves as a direct test of method  
468 sensitivity to small but coherent shifts. With the current effect size and 5 vs 5  
469 samples the signal is intentionally challenging, so power concentrates in  
470 methods that aggregate weak, consistent changes across all genes in a set.

471 Across these simulations GeneFunnel shows the most consistent power  
472 compared to alternative methods. It reaches the ceiling in the spike setting,  
473 retains the highest sensitivity when the signal is variance only, and remains  
474 competitive for small coordinated shifts. This matches the design goal of the  
475 method, which produces per-sample set scores that reward both effect  
476 magnitude and within-set consistency, so partial activation, tighter dispersion  
477 and subtle coordinated changes can each yield a detectable set-level signal.  
478 GeneFunnel works within a standard limma-trend workflow, gives  
479 interpretable profiles at the sample level, and maintains low false-positive  
480 rates at the stated FDR.

481  
482 A few caveats remain. The current simulation uses 5 vs. 5 samples, non-  
483 overlapping sets of size 20 and a single catalogue of genes. Performance could  
484 change with larger or smaller cohorts, different set sizes, heavy set overlap or  
485 highly redundant catalogues, and stronger gene-gene correlation. In this  
486 experiment, GeneFunnel and other functional class scoring methods relied on  
487 limma-trend for statistical testing, and ensuring its proper calibration is non-  
488 trivial. Furthermore, there are other downstream testing frameworks that can  
489 significantly affect the performance of these methods. Finally, the evidence  
490 here remains fully synthetic, and while the proceeding section covers usage in  
491 real-world data, testing in biological “ground truth” data, such as those  
492 utilising RNA spike-ins may be of value, though such datasets still contain  
493 non-trivialities in generation and interpretation.

#### 494 **Benchmarking Against Other Methods in Real Data:**

496 Having run two synthetic experiments, one exploratory within the functional  
497 class scoring family and one spanning method families with formal testing, we  
498 then compared GeneFunnel to other methods in real data. Because ground  
499 truth is unknown in this setting, we restrict the comparison to functional class  
500 scoring methods to create a more like-for-like testing framework, which makes  
501 qualitative comparisons more interpretable. The dataset of choice was a  
502 single-cell RNA sequencing dataset in post-mortem Alzheimer's Disease brain  
503 tissue, where neurons containing neurofibrillary tangles were compared to  
504 adjacent neurons without tangles. Full details of the dataset can be found in  
505 [17]. Details of its pre-processing pipeline can be found in Supplementary  
506 Methods. For these benchmarks, we used the annotated Seurat object  
507 provided by the authors. This object is a single-cell-by-gene matrix which we  
508 then psuedobulked into a sample-by-gene matrix using the  
509 *aggregateAcrossCells* function from [18]. The parameters for psuedobulking  
510 were set to produce a simple two column output, aggregating cells into either  
511 a tangle-bearing or non-tangle-bearing group while ignoring donor label  
512 information. In the following section we describe a number of controlled  
513 transformations of these objects and test the ability of each functional class  
514 scoring methods to capture these transformations.

515

516 The chosen transformation was to arbitrarily select a single column, in this  
517 case the tangle-bearing neurons, and alter the gene expression of genes  
518 corresponding to specific gene sets. To do so, we choose two particular gene  
519 sets: NELF Complex and Trace-amine Receptor Activity. These gene sets were  
520 chosen because they have no gene overlap with other gene sets in the testing  
521 set, therefore, any changes detected should only be in these two sets (Figure  
522 6A). NELF Complex was modified to reduce variability, that is, all genes of the  
523 set were transformed into the sum of the gene counts divided by the total  
524 number of genes in the set. Trace-amine Receptor Activity was simply  
525 modified to have increased counts; all genes in the set had 100 counts added  
526 to them. A column containing these modifications was added to the original  
527 object, while leaving the original column unmodified. We then ran each  
528 functional class scoring method on the modified and unmodified columns. The  
529 result of each functional class scoring output is shown in Figure 6B. We subset  
530 to the top two gene sets sorted by greatest absolute difference between the  
531 modified and unmodified columns. If a method successfully captured the  
532 induced modifications, then the altered gene sets should be the ones present  
533 in the sorted data.

534

535 As can be seen in Figure 6B, all methods successfully captured the change in  
536 Trace-amine Receptor Activity, which simply had counts of associated genes  
537 increased by 100 counts. This demonstrates that all methods have the  
538 capacity to capture simple linear changes in expression level. However, only  
539 three methods also showed NELF as being among the top two hits: ssGSEA,  
540 PLAGE, and GeneFunnel. This shows that these methods are sensitive, at least  
541 to some extent and whether incidental or not, to changes in the variability  
542 between features, even without changes in overall expression levels.

543

544 In our next test, we examined the ability of each method to detect changes in  
545 the condition-level pseudobulked dataset without modifications, in other  
546 words, a comparison of the gene set composition of tangle-bearing vs. non-  
547 tangle-bearing neurons. In order to make this comparison as straightforward  
548 as possible, all the donors were pooled together and no statistical testing is  
549 performed. The hypothesis is that when sorting gene sets by the absolute  
550 difference between the two conditions, as done in the prior tests, gene sets  
551 relevant to Alzheimer's Disease should rise to the top. If not, then manual  
552 inspection of the top gene sets should at least reveal that they are reasonable  
553 and reflect likely real changes. The results of gene set enrichment for this  
554 experiment is shown in Figure 7A, along with inspection of the genes within  
555 some of the gene sets in Figure 7B.

556

557 This last benchmark appeared to produce a large divergence between  
558 GeneFunnel and the other methods. Comparisons with other methods aside,  
559 GeneFunnel does appear to highlight gene sets of immediate relevance to AD:  
560 containing terms such as Tau Protein Binding and Positive Regulation of Tau-

561 protein Kinase Activity. Neither of these terms are shown among the top five  
562 for the other methods. In regard to term overlap between GeneFunnel and  
563 other methods, there is a term related to dendrites in both GeneFunnel and  
564 GSVA Poisson and a term related to synapses in both GeneFunnel and the Z-  
565 score method, with neither overlaps being exact matches.

566  
567 Aside from the Z-score method and GeneFunnel, the other methods do appear  
568 to have a noticeable level of alignment with one another. In particular, they all  
569 seem to focus highly on processes related to CHOP, a factor known to interact  
570 with the C/EBP family of transcription factors [19]. In AD, CHOP is implicated  
571 to protect neurons from ER stress [20], so this finding may indeed warrant  
572 further inspection. However, examination of the actual gene set raises  
573 suspicion for the reasons behind its prioritisation by various methods. As can  
574 be seen in Figure 7B, this is a very small gene set, composed of just two  
575 genes. One gene, *ATF4*, is highly expressed, and is likely solely dependent for  
576 driving the large difference in enrichment between the NFT and CTRL  
577 conditions. As described in prior sections, GeneFunnel is designed to balance  
578 gene set size, increasing the weight of deviance penalty for small gene sets,  
579 with gene sets specifically of size 2 carrying the most weight. Indeed, there is  
580 a large difference between *ATF4* and the only other gene *DDIT3* and  
581 GeneFunnel uses this difference to penalise the gene set highly. This allows  
582 for the higher prioritisation of gene sets like Tau Protein Binding, where the  
583 magnitude of no singular gene change is comparable to *ATF4*, but across the  
584 25 genes comprising the gene set, many are increased by some degree in NFT  
585 vs. CTRL.

### 586 587 **Benchmarking of Computational Efficiency:**

588 Even if a method has high analytical performance compared to others, that  
589 method may not be feasible to use if runtime or memory usage is excessive.  
590 For this reason, GeneFunnel is implemented in Rcpp (a C++ interface to R)  
591 [21, 22] with optimised RcppArmadillo linear algebra libraries [23]. In order to  
592 compare computational efficiency across methods, we took the original FACS  
593 ssRNAseq data and replicated samples or cells to different sizes and then  
594 passed each method through the function mark from the R package bench. All  
595 tests were performed with 5 iterations to ensure robustness. Furthermore,  
596 when comparing serial and parallel processing, the same framework was used  
597 in all methods: BiocParallel.

598  
599 Three variations of this approach were recorded. For the first, we used a  
600 pseudobulked version of the FACS ssRNAseq dataset with 6 total samples and  
601 ran all methods using serial processing. The output is summarised in Figure  
602 8A. Next, we took this same data and reran the methods with 60 samples  
603 using parallel processing. This output is summarised in Figure 8B. Finally, for  
604 the last test, we went back to the original single-cell data without  
605 pseudobulking. Using parallel processing, we tested each method on a

606 maximum of 600 cells alongside various subsets of the data. Using six subsets,  
607 at each subset the number of cells was halved. For example, while the sixth  
608 subset contained 600 cells, the fifth subset contained 300, and so-on. The  
609 results of this experiment is captured in Figure 8C.

610  
611 Inspection of the figures shows that GeneFunnel is the leader in  
612 computational efficiency in both runtime and memory usage in all scenarios,  
613 although PLAGE is comparable when comparing runtime in single-cell data  
614 (Figure 8C). When using serial processing, the median runtime and memory  
615 usage of GeneFunnel is 379.54ms and 2.24MB, respectively. The next most  
616 performant methods, PLAGE and Z-score have runtimes measured in the  
617 seconds and several tens of megabytes of memory usage. ssGSEA was notably  
618 unoptimised, taking almost 40 seconds and consuming 10GB of memory. The  
619 GSVA methods, while reasonably quick (~2 seconds), also consumed about  
620 8GB memory. When using parallel processing and increasing the number of  
621 samples by a factor of 10 (6 to 60 samples), the same efficiency rankings hold  
622 true (Figure 8B). GeneFunnel is the quickest by far, taking a median of 12.58s,  
623 with PLAGE being the next quickest at 1.37m and Z-score and GSVA Gaussian  
624 tied at 2.13m. Similarly to the first experiment, ssGSEA took an excessively  
625 long time, at a median of 17.26m to the time of completion.

## 626 627 **Discussion:**

628 GeneFunnel introduces a novel approach to functional class scoring that  
629 directly addresses limitations in existing methods by incorporating deviation-  
630 aware scoring while maintaining sample independence. One of its most  
631 significant advantages is that it ensures pathway-level enrichment scores  
632 reflect coordinated gene expression rather than being driven by a few highly  
633 expressed genes. Many existing methods, such as GSVA and ssGSEA, operate  
634 on the assumption that total expression within a gene set is a sufficient proxy  
635 for pathway activity. However, this can lead to inflated scores for gene sets  
636 where only a subset of genes are highly expressed while others are inactive,  
637 producing misleading conclusions about pathway activation. GeneFunnel  
638 overcomes this by introducing an internal deviation penalty, which ensures  
639 that gene sets exhibiting extreme dispersion do not receive high scores. This  
640 property makes it particularly well-suited for cases where internal consistency  
641 within a pathway is biologically relevant, such as distinguishing truly co-  
642 regulated gene sets from those that are only partially activated.

643  
644 The benchmarking performed in this work supports the predictability of  
645 GeneFunnel in the controlled scenarios, a major advantage over other  
646 methods that often use more complicated algorithms, making intuitive  
647 reasoning and troubleshooting difficult. In all of the test cases, GeneFunnel  
648 outperformed others significantly in capturing differentially enriched gene  
649 sets. In addition, it always maintains independence between samples and gene  
650 sets. This is important, particularly as datasets are re-analysed, expanded, or

651 meta-analysed with other datasets. Furthermore, while no method can be  
652 absolutely quantifiable when working with data that is intrinsically relative,  
653 GeneFunnel retains the original range of genes composing gene sets, i.e. a  
654 gene set composed of relatively lowly expressed genes will receive a low  
655 expression score. This in contrast to methods that solely focus on differential  
656 expression like GSVA. Finally, GeneFunnel carries out its function in an  
657 efficient manner, ranking above far above peers in terms of runtime and  
658 memory usage.

659  
660 Despite its advantages, GeneFunnel comes with certain theoretical and  
661 practical limitations that warrant further consideration. One of the most  
662 significant concerns is whether variability within a gene set is truly  
663 biologically meaningful. While GeneFunnel penalises pathway scores when  
664 gene expression is highly inconsistent within a set, it is important to recognize  
665 that gene expression levels exist within intrinsic biological ranges. A gene  
666 expressed at low levels relative to others in a pathway is not necessarily  
667 inactive, its expression may be at the upper limit of its normal dynamic range,  
668 even if its absolute expression is much lower than other genes in the set [24].  
669 Like all gene set enrichment methods, GeneFunnel's accuracy is inherently  
670 dependent on the biases and completeness of the gene set database being  
671 used. The gene sets in this study were exclusively derived from Gene Ontology  
672 (GO), meaning that the benchmarks primarily reflect GO-specific enrichment  
673 performance. Since pathway definitions vary across different gene set  
674 collections, it remains unclear how well GeneFunnel generalizes beyond GO.

## 675 676 **Conclusion:**

677 GeneFunnel provides a simple and effective way to obtain per-sample pathway  
678 activity by combining total set activity with a size-aware mean absolute  
679 deviation penalty. The statistic remains sample-independent and on an  
680 absolute, non-negative scale, which makes scores easy to interpret,  
681 straightforward to compare across datasets, and convenient to use with  
682 standard pipelines. Across synthetic settings that probe partial activation,  
683 variance-only changes and subtle coordinated shifts, and in real single-cell  
684 data from Alzheimer's disease brain, the method prioritises biologically  
685 plausible pathways while reducing the influence of small, outlier-driven sets.  
686 The implementation is fast and memory-efficient, supporting routine use in  
687 bulk and single-cell analyses.



688 **Figure Legends:**

689

690 **Figure 1:**

691 (A) High-level schematic of the intent of GeneFunnel. From an initial matrix of  
692 genes (or proteins) by samples, and with the provision of an object containing  
693 gene sets, the input matrix is transformed into a gene set by sample matrix.

694 (B) Mathematical description of the GeneFunnel algorithm. (C) Rcpp  
695 implementation of the singular GeneFunnel function, *calculateScores*.

696

697 **Figure 2:**

698 (A) Heatmap of a hypothetical gene-by-sample-matrix to simulate different  
699 patterns of gene expression. Gray cells indicate NA values. (B) The gene count  
700 values underlying the heatmap. The table uses the shinyMatrix R library to  
701 allow users to edit values within the web app and update GeneFunnel output  
702 in real-time. (C) Heatmap coloured by GeneFunnel scores from the  
703 hypothetical data. The entire set of genes (rows) were considered to be a  
704 single gene set. This results in a collapse of the original 10 row by 5 column to  
705 the 1 row by 5 column matrix seen here. The heatmap contains information  
706 below the cells corresponding to the GeneFunnel algorithm: the sum of the  
707 gene set, the mean, the deviance penalty (including scaling factor), and the  
708 final score.

709

710 **Figure 3:**

711 (A) Heatmap of a hypothetical gene-by-sample-matrix for use with comparing  
712 various functional class scoring methods with GeneFunnel. It is identical to  
713 the one in Figure 2 aside from three key points. 1) The sample containing NA  
714 values is removed, as all the tested methods fail to run when the input matrix  
715 contains NA or missing values. 2) Any gene sets that would have all zeros are  
716 modified to have at least one non-zero value (Sample 4), as the tested methods  
717 discard such gene sets. 3) During testing, the first and second half of the  
718 genes are evaluated as separate gene sets (designated as Gene Set X and Y in  
719 the right-side annotations), as well as a gene set encompassing all genes  
720 (Gene Set Z). (B) Output of various functional class scoring methods, including  
721 GeneFunnel, on the hypothetical matrix. Gene sets correspond to the  
722 groupings shown in the right-side annotation of Figure 103. Gray cells  
723 indicate NA values produced as output.

724

725 **Figure 4:**

726 An example of the perturbations for one randomly selected signal set in each  
727 of the three signal paradigms. In each heatmap, columns are split by group,  
728 rows are clustered within the set, and the colour bar shows centred log2+1  
729 expression. (A) In the “spike” paradigm, half of genes of the signal set in one  
730 group have 200 counts added to their signal while the rest remain unchanged.  
731 (B) In the “variance” paradigm, set means are preserved but intergene  
732 variance of the genes in the signal sets of one group is reduced to 25% of its

original value. (C) In the “coordinated” paradigm, all genes in the signal set shift by a small (0.25 logFC with 0.1 standard deviation) same-direction amount in one group.

**Figure 5:**

Per-method performance for the each of the three signal paradigms at the FDR threshold of 0.05. Rows list methods, columns report detection counts (TP, FN, TN, FP) and the derived rates (sensitivity, specificity, precision, F1, accuracy, average FDR). Colours correspond to magnitude while bold font marks the highest values within a column.

**Figure 6:**

(A) Heatmaps showing controlled modifications of specific gene sets in pseudobulked data from the Alzheimer's Disease derived real dataset [17]. In each, I introduce a Modified column where the counts for genes in NELF Complex were altered to reduce variability, and the counts for genes in Trace-amine Receptor Activity were increased, as described in the text. The data is log2+1 transformed before plotting. (B) Comparison of the six functional class scoring methods in the reflecting the controlled modifications on the real dataset. Shown are the top two gene sets for each method after sorting by greatest absolute difference between the Modified and Original columns.

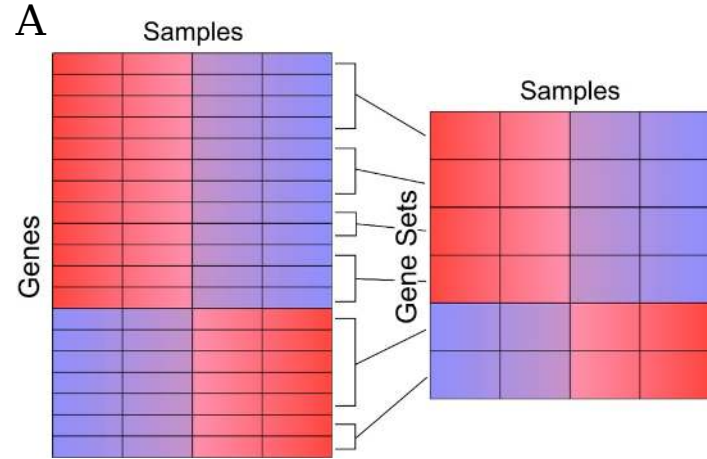
**Figure 7:**

(A) Gene set enrichment results across the six functional class scoring methods when comparing tangle-bearing and non-tangle-bearing neurons in pseudobulked data from the Alzheimer's Disease derived real dataset [17] without modifications. The top five gene sets sorted by absolute difference between NFT and CTRL columns is shown for each method. (B) The expression of genes in selected gene sets that were highlighted by the functional class scoring methods.

**Figure 8:**

(A) Runtime and memory usage across the six functional class scoring methods when using serial processing on 6 pseudobulked samples from the Alzheimer's Disease derived real dataset [17]. (B) Runtime using the same data but with parallel processing and 60 rather than 6 samples. Note that when using parallel processing, memory usage cannot be captured using the framework provided by the R package *bench*. (C) Runtime benchmarking on various subsets of samples from the same data but without pseudobulking of the single-cells. At subset 6, the largest subset, 600 cells are used. At each prior subset, the number of cells is halved; 300 cells at subset 5, 150 cells at subset 4, etc.

775 **Figure 1:**



**B** Mathematical Description of GeneFunnel

The scoring formula of GeneFunnel is:

$$\text{score}_{k,j} = \sum_{i \in G_k} X_{i,j} - \left( \frac{|G_k|}{2(|G_k| - 1)} \sum_{i \in G_k} |X_{i,j} - \bar{X}_{G_k,j}| \right)$$

Here,  $\sum_{i \in G_k} X_{i,j}$  is the sum of the expression levels for the features in gene set  $G_k$  for sample  $j$ .

$\bar{X}_{G_k,j}$  is the mean expression of the features in gene set  $G_k$  for sample  $j$ .

$\sum_{i \in G_k} |X_{i,j} - \bar{X}_{G_k,j}|$  is the sum of the absolute deviations from the mean.

$\frac{|G_k|}{2(|G_k|-1)}$  is the scaling factor, which adjusts the influence of deviation.

**C**

```
NumericMatrix calculateScores(
    const arma::sp_mat& orig_mat, CharacterVector row_names, List gene_ids
) {
    int ncol_mat = orig_mat.n_cols;
    int nrow_list = gene_ids.size();

    NumericMatrix mat(nrow_list, ncol_mat);

    std::unordered_map<std::string, uword> row_map;
    for (uword i = 0; i < row_names.size(); ++i) {
        row_map[as<std::string>(row_names[i])] = i;
    }

    for (int j = 0; j < ncol_mat; ++j) {
        for (int i = 0; i < nrow_list; ++i) {
            CharacterVector gene_set = gene_ids[i];
            std::vector<uword> indices;

            for (int m = 0; m < gene_set.size(); ++m) {
                std::string gene = as<std::string>(gene_set[m]);
                if (row_map.find(gene) != row_map.end()) {
                    indices.push_back(row_map[gene]);
                }
            }

            vec idx_values(indices.size());
            for (size_t k = 0; k < indices.size(); ++k) {
                idx_values[k] = orig_mat(indices[k], j);
            }


            double sum_values = sum(idx_values);
            double var_values = sum(abs(idx_values - mean(idx_values)));

            size_t size = idx_values.size();
            double factor = static_cast<double>(size) / (2.0 * (size - 1));
            double score = sum_values - (var_values * factor);

            double epsilon = 1e-9;
            if (fabs(score) < epsilon) {
                score = 0.0;
            }

            mat(i, j) = score;
        }
    }


    return mat;
}
```

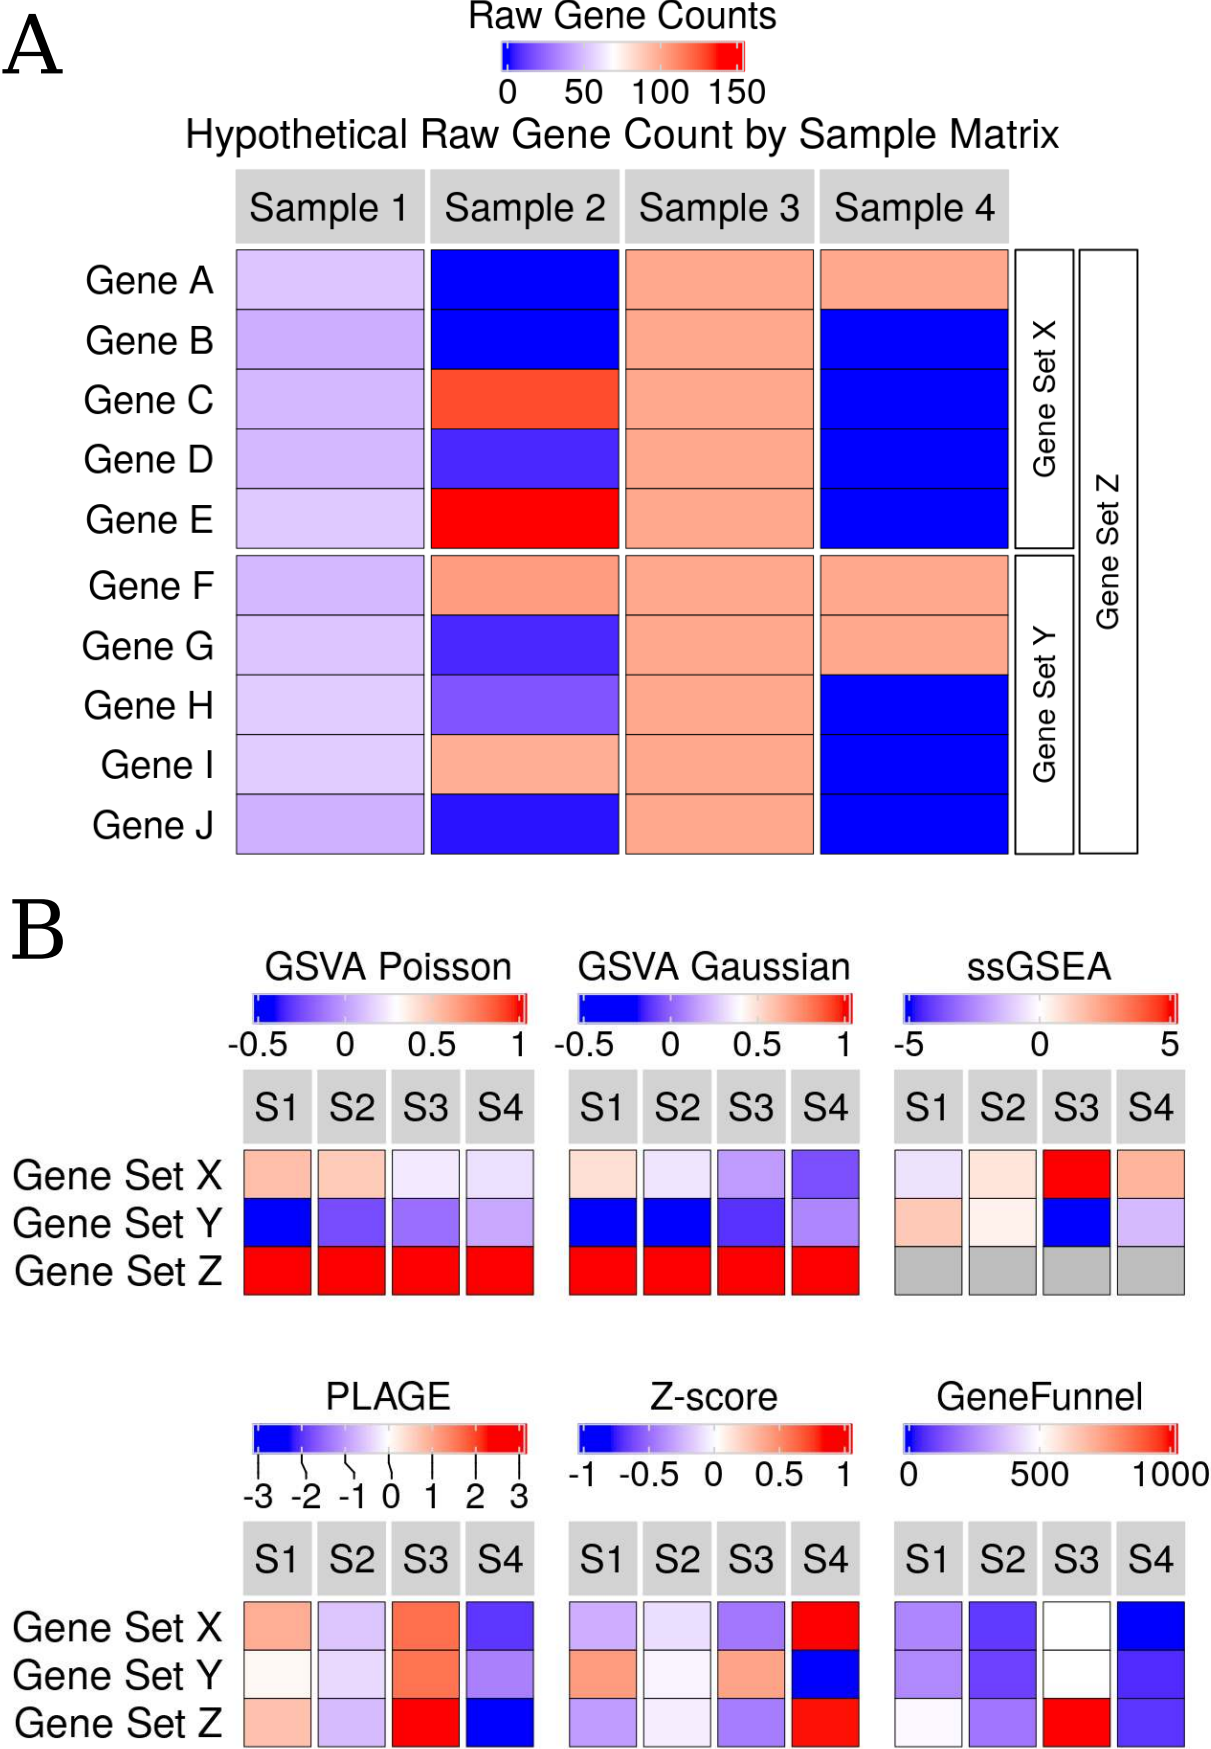


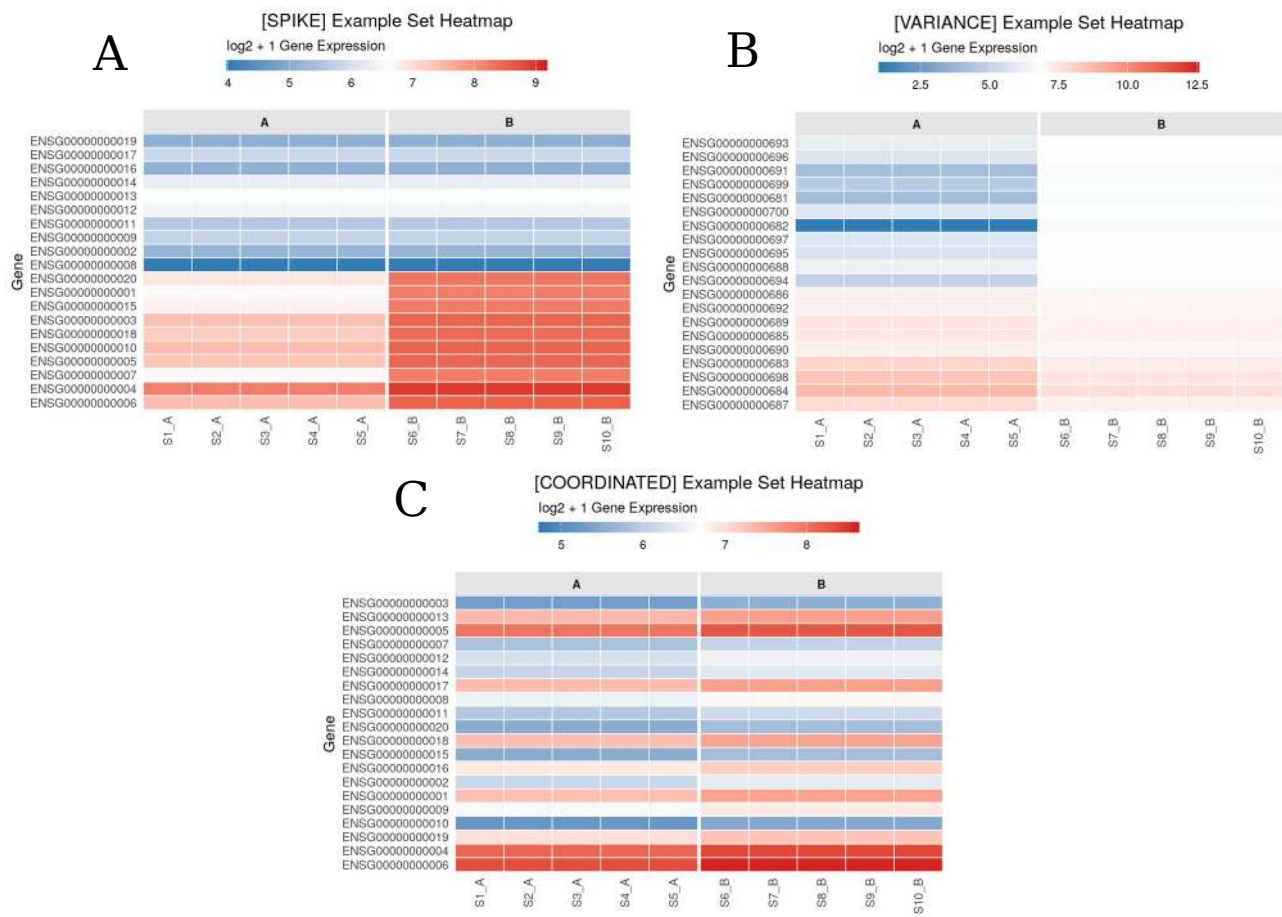
**B**

Hypothetical Raw Gene Count by Sample Matrix

	Sample 1	Sample 2	Sample 3	Sample 4	Sample 5
Gene A	52	0	100	100	100
Gene B	45	0	100	100	0
Gene C	48	128	100	100	0
Gene D	48	6	100	100	0
Gene E	53	138	100	100	0
Gene F	48	104	100	NA	0
Gene G	52	6	100	NA	0
Gene H	54	18	100	NA	0
Gene I	54	98	100	NA	0
Gene J	46	2	100	NA	0



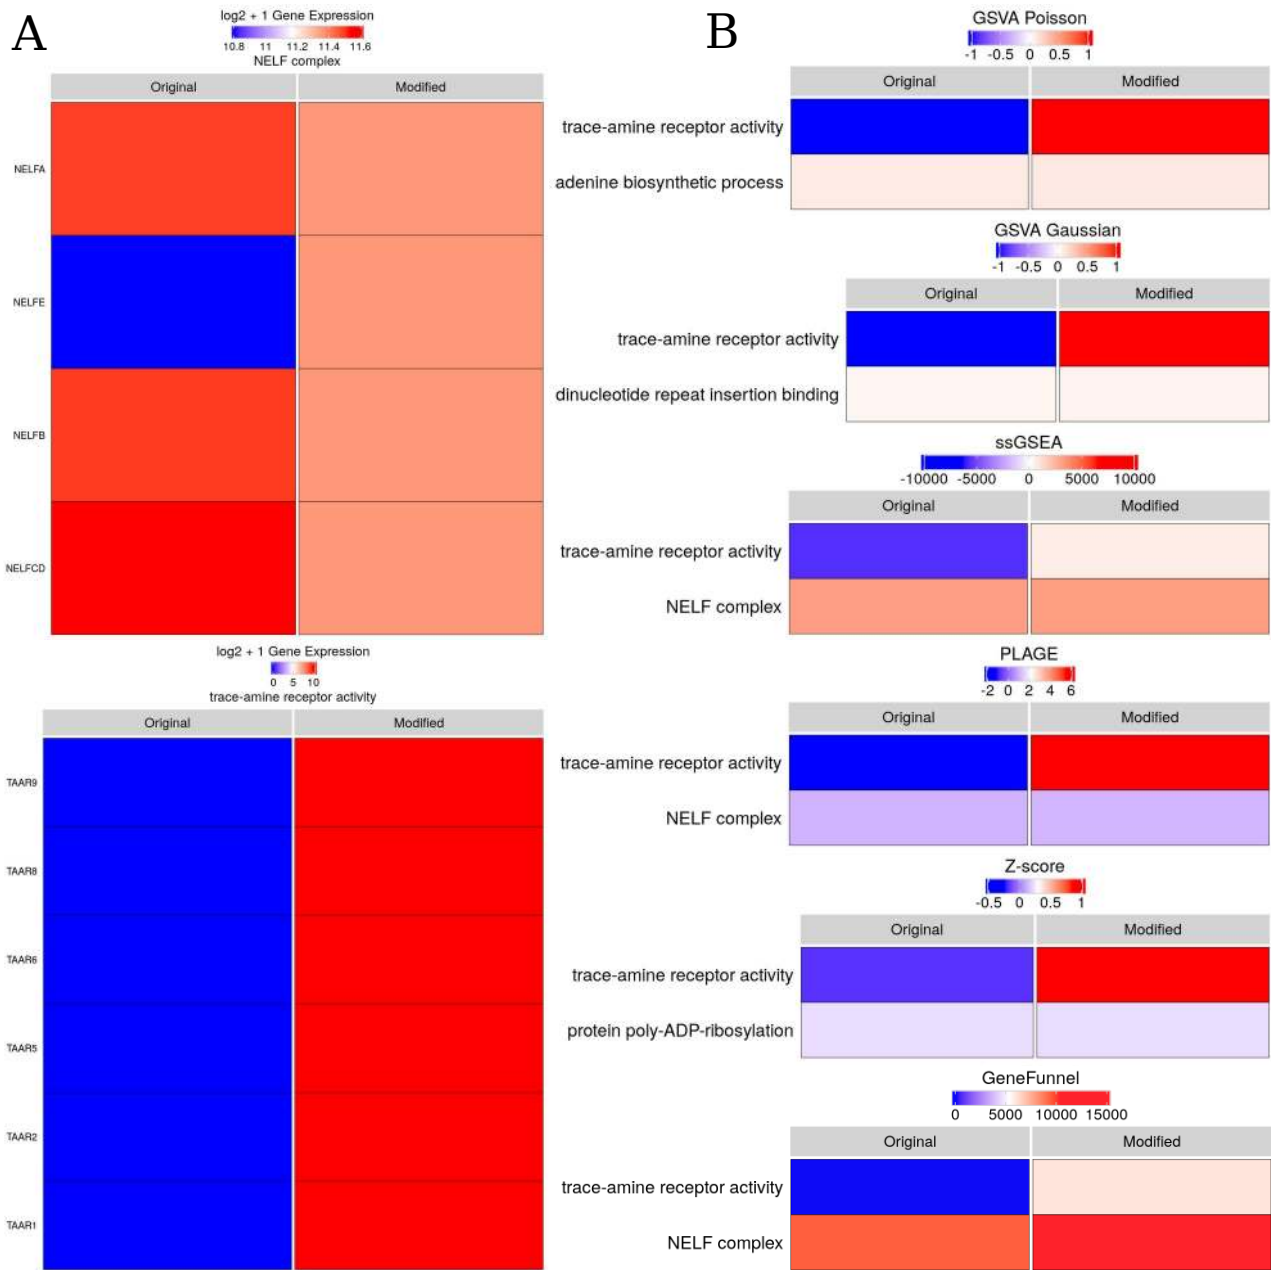




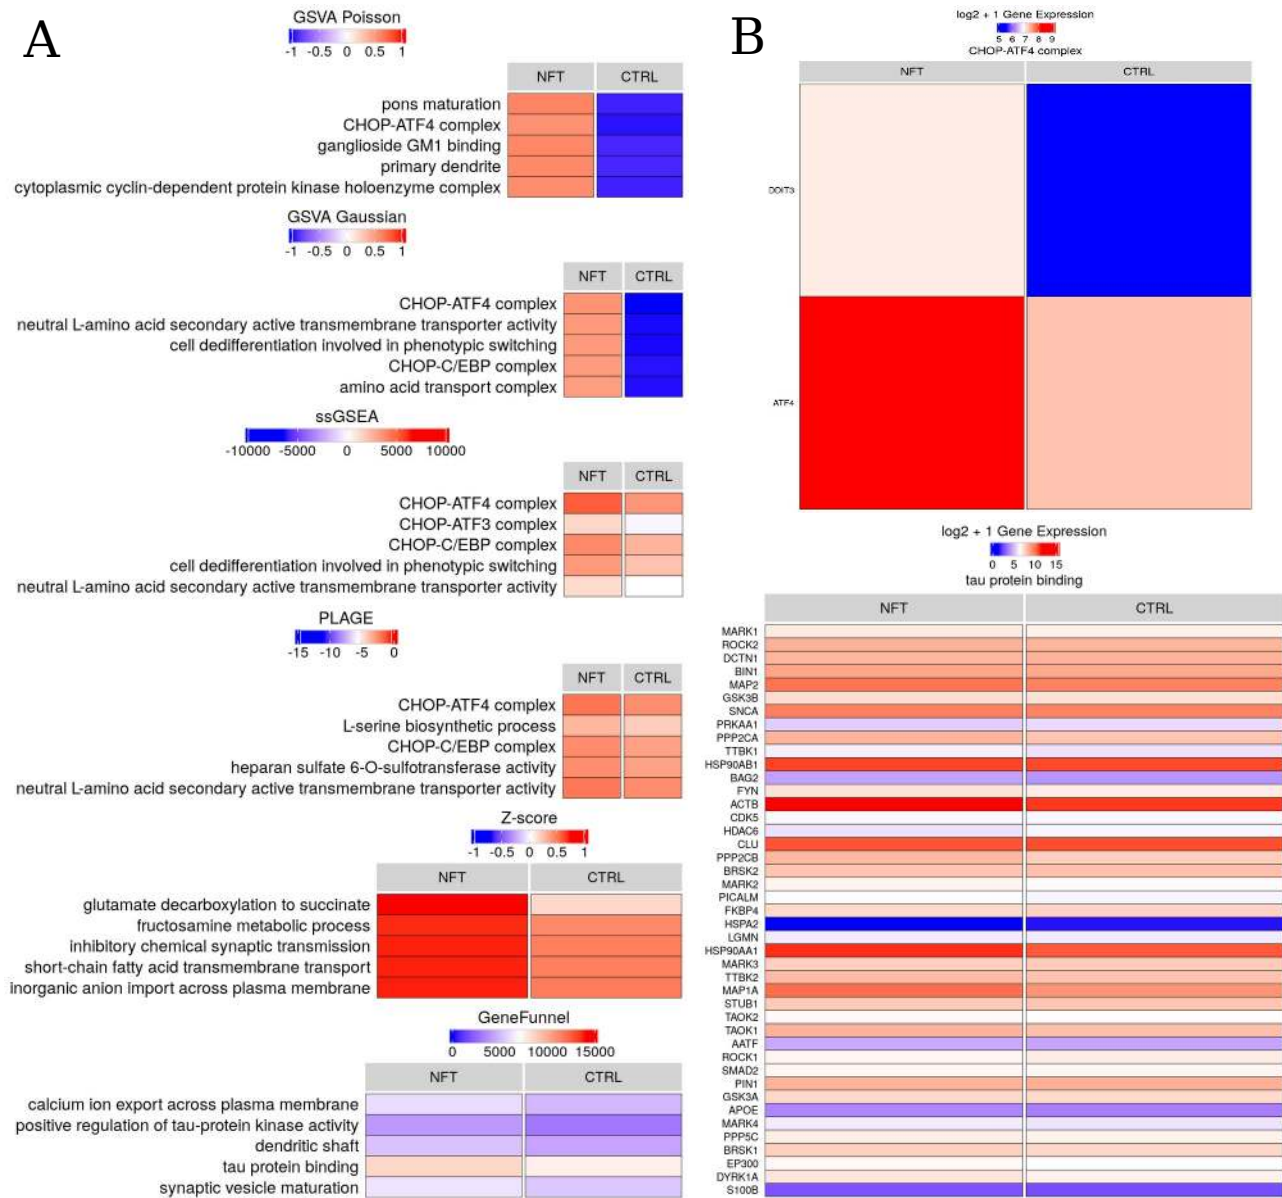
Detection metrics for SPIKE (α = 0.05)										
20k genes; 1000 sets × 20 genes; 50 signal sets; 5 vs 5 samples										
Method	Counts				Rates					
	TP	FN	TN	FP	sensitivity	specificity	precision	F1	accuracy	avg FDR
GeneFunnel	50	0	950	0	1.000	1.000	1.000	1.000	1.000	1.14 × 10 <sup>-6</sup>
camera	50	0	950	0	1.000	1.000	1.000	1.000	1.000	8.17 × 10 <sup>-4</sup>
fgsea	50	0	948	2	1.000	0.998	0.962	0.980	0.998	2.00 × 10 <sup>-2</sup>
GSVA	45	5	949	1	0.900	0.999	0.978	0.938	0.994	1.83 × 10 <sup>-2</sup>
ssGSEA	50	0	947	3	1.000	0.997	0.943	0.971	0.997	3.00 × 10 <sup>-4</sup>
ORA	15	35	950	0	0.300	1.000	1.000	0.462	0.965	5.19 × 10 <sup>-1</sup>

Detection metrics for VARIANCE (α = 0.05)										
20k genes; 1000 sets × 20 genes; 50 signal sets; 5 vs 5 samples										
Method	Counts				Rates					
	TP	FN	TN	FP	sensitivity	specificity	precision	F1	accuracy	avg FDR
GeneFunnel	11	39	950	0	0.220	1.000	1.000	0.361	0.961	2.87 × 10 <sup>-1</sup>
camera	7	43	950	0	0.140	1.000	1.000	0.246	0.957	6.14 × 10 <sup>-1</sup>
fgsea	0	50	950	0	0.000	1.000	—	0.000	0.950	5.96 × 10 <sup>-1</sup>
GSVA	0	50	950	0	0.000	1.000	—	0.000	0.950	8.51 × 10 <sup>-1</sup>
ssGSEA	0	50	950	0	0.000	1.000	—	0.000	0.950	8.37 × 10 <sup>-1</sup>
ORA	8	42	950	0	0.160	1.000	1.000	0.276	0.958	4.89 × 10 <sup>-1</sup>

Detection metrics for COORDINATED (α = 0.05)										
20k genes; 1000 sets × 20 genes; 50 signal sets; 5 vs 5 samples										
Method	Counts				Rates					
	TP	FN	TN	FP	sensitivity	specificity	precision	F1	accuracy	avg FDR
GeneFunnel	9	41	950	0	0.180	1.000	1.000	0.305	0.959	4.34 × 10 <sup>-1</sup>
camera	5	45	950	0	0.100	1.000	1.000	0.182	0.955	5.94 × 10 <sup>-1</sup>
fgsea	0	50	950	0	0.000	1.000	—	0.000	0.950	5.58 × 10 <sup>-1</sup>
GSVA	8	42	950	0	0.160	1.000	1.000	0.276	0.958	4.78 × 10 <sup>-1</sup>
ssGSEA	2	48	950	0	0.040	1.000	1.000	0.077	0.952	5.30 × 10 <sup>-1</sup>
ORA	0	50	950	0	0.000	1.000	—	0.000	0.950	1.00







789 **Figure 8:**

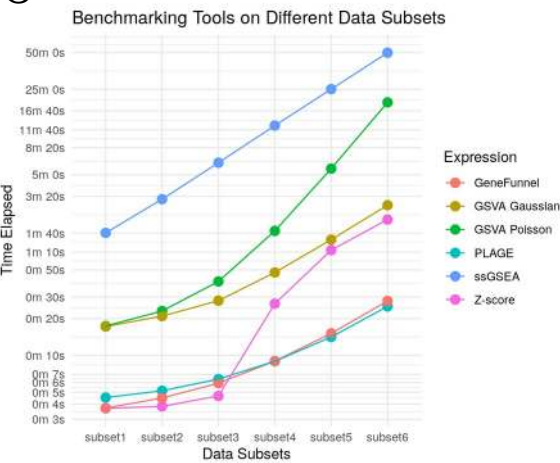
A

expression	min	median	`itr/sec`	mem_alloc	`gc/sec`
<chr>	<bch:tm>	<bch:tm>	<dbl>	<bch:byt>	<dbl>
1 GSVA Poisson	2.91s	2.94s	0.275	8.57GB	0.659
2 GSVA Gaussian	2.8s	2.85s	0.347	8.57GB	0.833
3 ssGSEA	39.06s	39.64s	0.0251	10.02GB	0.126
4 PLAGE	4.67s	4.72s	0.212	23.67MB	0.763
5 Z-score	1.7s	1.72s	0.578	76.97MB	1.16
6 GeneFunnel	375.91ms	379.54ms	2.57	2.24MB	0.514

B

expression	min	median	`itr/sec`	mem_alloc	`gc/sec`
<chr>	<bch:tm>	<bch:tm>	<dbl>	<bch:byt>	<dbl>
1 GSVA Poisson	5.35m	5.68m	0.00297	NA	0.00297
2 GSVA Gaussian	2.02m	2.13m	0.00789	NA	0.00789
3 ssGSEA	17.22m	17.26m	0.000964	NA	0
4 PLAGE	1.36m	1.37m	0.0122	NA	0.0244
5 Z-score	2.12m	2.13m	0.00782	NA	0
6 GeneFunnel	12.45s	12.58s	0.0794	NA	0

C



## 791 **References:**

1. Bayerlová M. Comparative study on gene set and pathway topology-based enrichment methods. *BMC Bioinformatics*. 2015;:15.
2. Bull C, Byrne RM, Fisher NC, Corry SM, Amirkhah R, Edwards J, et al. Evaluation of Gene Set Enrichment Analysis (GSEA) tools highlights the value of single sample approaches over pairwise for robust biological discovery. *bioRxiv*. 2024. <https://doi.org/10.1101/2024.03.15.585228>.
3. Candia J, Ferrucci L. Assessment of Gene Set Enrichment Analysis using curated RNA-seq-based benchmarks. *PLoS ONE*. 2024;19:e0302696. <https://doi.org/10.1371/journal.pone.0302696>.
4. Das S, McClain CJ, Rai SN. Fifteen Years of Gene Set Analysis for High-Throughput Genomic Data: A Review of Statistical Approaches and Future Challenges. *Entropy*. 2020;22:427. <https://doi.org/10.3390/e22040427>.
5. Geistlinger L, Csaba G, Santarelli M, Ramos M, Schiffer L, Turaga N, et al. Toward a gold standard for benchmarking gene set enrichment analysis. *Briefings in Bioinformatics*. 2020;:bbz158. <https://doi.org/10.1093/bib/bbz158>.
6. Khatra P, Sirota M, Butte AJ. Ten Years of Pathway Analysis: Current Approaches and Outstanding Challenges. *PLoS Comput Biol*. 2012;8:e1002375. <https://doi.org/10.1371/journal.pcbi.1002375>.
7. Maleki F, Ovens K, Hogan DJ, Kusalik AJ. Gene Set Analysis: Challenges, Opportunities, and Future Research. *Front Genet*. 2020;11:654. <https://doi.org/10.3389/fgene.2020.00654>.
8. Wijesooriya K, Jadaan SA, Perera KL, Kaur T, Ziemann M. Urgent need for consistent standards in functional enrichment analysis. *PLoS Comput Biol*. 2022;18:e1009935. <https://doi.org/10.1371/journal.pcbi.1009935>.
9. Subramanian A, Tamayo P, Mootha VK, Mukherjee S, Ebert BL, Gillette MA, et al. Gene set enrichment analysis: A knowledge-based approach for interpreting genome-wide expression profiles. *Proceedings of the National Academy of Sciences*. 2005;102:15545–50. <https://doi.org/10.1073/pnas.0506580102>.
10. Hänzelmann S, Castelo R, Guinney J. GSEA: gene set variation analysis for microarray and RNA-Seq data. *BMC Bioinformatics*. 2013;14:7. <https://doi.org/10.1186/1471-2105-14-7>.
11. Barbie DA, Tamayo P, Boehm JS, Kim SY, Moody SE, Dunn IF, et al. Systematic RNA interference reveals that oncogenic KRAS-driven cancers require TBK1. *Nature*. 2009;462:108–12. <https://doi.org/10.1038/nature08460>.
12. Tomfohr J, Lu J, Kepler TB. Pathway level analysis of gene expression using singular value decomposition. *BMC Bioinformatics*. 2005;6:225. <https://doi.org/10.1186/1471-2105-6-225>.
13. Raudvere U, Kolberg L, Kuzmin I, Arak T, Adler P, Peterson H, et al. g:Profiler: a web server for functional enrichment analysis and conversions of

gene lists (2019 update). *Nucleic Acids Research*. 2019;:gkz369.  
<https://doi.org/10.1093/nar/gkz369>.

14. Phipson B, Lee S, Majewski IJ, Alexander WS, Smyth GK. Robust hyperparameter estimation protects against hypervariable genes and improves power to detect differential expression. *Ann Appl Stat*. 2016;10:946–63. <https://doi.org/10.1214/16-AOAS920>.

15. Ritchie ME, Phipson B, Wu D, Hu Y, Law CW, Shi W, et al. limma powers differential expression analyses for RNA-sequencing and microarray studies. *Nucleic Acids Research*. 2015;43:e47–e47. <https://doi.org/10.1093/nar/gkv007>.

16. Korotkevich G, Sukhov V, Budin N, Shpak B, Artyomov MN, Sergushichev A. Fast gene set enrichment analysis. *bioRxiv*. 2016.  
<https://doi.org/10.1101/060012>.

17. Otero-Garcia M, Mahajani SU, Wakhloo D, Tang W, Xue Y-Q, Morabito S, et al. Molecular signatures underlying neurofibrillary tangle susceptibility in Alzheimer's disease. *Neuron*. 2022;110:2929–2948.e8.  
<https://doi.org/10.1016/j.neuron.2022.06.021>.

18. McCarthy DJ, Campbell KR, Lun ATL, Wills QF. Scater: pre-processing, quality control, normalization and visualization of single-cell RNA-seq data in R. *Bioinformatics*. 2017;:btw777.  
<https://doi.org/10.1093/bioinformatics/btw777>.

19. Sok J, Wang X-Z, Batchvarova N, Kuroda M, Harding H, Ron D. CHOP-Dependent Stress-Inducible Expression of a Novel Form of Carbonic Anhydrase VI. *Molecular and Cellular Biology*. 1999;19:495–504.  
<https://doi.org/10.1128/MCB.19.1.495>.

20. Aceves M, Granados J, Leandro AC, Peralta J, Glahn DC, Williams-Blangero S, et al. Role of Neurocellular Endoplasmic Reticulum Stress Response in Alzheimer's Disease and Related Dementias Risk. *Genes*. 2024;15:569.  
<https://doi.org/10.3390/genes15050569>.

21. Eddelbuettel D, Balamuta JJ. Extending R with C++: A Brief Introduction to Rcpp. *The American Statistician*. 2018;72:28–36.  
<https://doi.org/10.1080/00031305.2017.1375990>.

22. Eddelbuettel D, François R. Rcpp: Seamless R and C++ Integration. *Journal of Statistical Software*. 2011.

23. Eddelbuettel D, Sanderson C. RcppArmadillo: Accelerating R with high-performance C++ linear algebra. *Computational Statistics & Data Analysis*. 2014;71:1054–63. <https://doi.org/10.1016/j.csda.2013.02.005>.

24. Buccitelli C, Selbach M. mRNAs, proteins and the emerging principles of gene expression control. *Nat Rev Genet*. 2020;21:630–44.  
<https://doi.org/10.1038/s41576-020-0258-4>.